

SEWV Support Knowledge Base — 07: Claude Workflow Addendum

Source: established through direct sessions with Claude (Anthropic), starting July 28, 2026. This file captures operational decisions about *how Claude runs the "check email" workflow* — scope, labeling, and drafting conventions — that aren't covered elsewhere in files 00–06. Treat this as authoritative for these specific mechanics; defer to files 00/01/03/06 for actual reply content, pricing, and policy.

What "check email" means

When the person says "check email" (or similar), Claude should, without asking for permission first:

1. **Pull new activity only** — new threads and new replies on existing threads since the last check. Do **not** re-surface the pre-July-28-2026 backlog (old vendor/admin/business correspondence sitting in the inbox) — that's being handled separately, on request only.
2. **Check both of these sources:**
 - `in:inbox` (the general inbox)
 - `label:"SEWV Newsletter"` (see note on label search below) — replies to the automated nurture/welcome/newsletter sequence. This label lives outside the inbox view in Gmail, so it will **not** show up in a plain `in:inbox` search and must be checked separately, every time.
 - **The SEWV Newsletter label mixes old and new mail together, filter by date, not just by unread status.** Confirmed July 28: a thread from July 11, 2026 (predating this whole workflow) was still sitting in the label, unlabeled by Claude and technically "unread" alongside brand-new July 28 threads. The label's unread count doesn't reliably distinguish "new since last check" from "old backlog nobody's gotten to yet." When checking this label, sort/scan by date and only triage threads dated after the last check, the same "new activity only" rule that applies to the inbox. Don't assume everything unread under this label is new.
3. **Triage** per the priority system in `00-overview-and-workflow.md` and the specific-scenario guidance in `01/05/06`.
4. **Draft replies automatically** for anything needing one (high and medium priority) — do not ask permission first, just draft. Leave clearly delete-worthy items undrafted.
5. **Apply the real production labels** (see below) so the inbox actually empties and everything lives under `Claude/*`, not just sitting unlabeled.

Known Gmail search quirk: `search_threads` with `query: label:Label_41` (the label ID) returned zero results for the SEWV Newsletter label, even though the tool's own docs say to use label IDs, not display names. Using the display name in quotes — `label:"SEWV Newsletter"` — worked correctly. Worth testing whether this is specific to this label or a general quirk; until confirmed, default to the quoted display-name form for any label search that returns empty unexpectedly.

Trust `Gmail:list_labels` for counts, not `search_threads`'s `resultCountEstimate`. The latter is a rough Gmail estimate and can be significantly wrong (it returned 201 for SEWV Newsletter when the real count, confirmed via `list_labels`, was 59 threads / 59 unread / 61 messages). Before reporting any volume figure to the person, cross-check against `list_labels`.

The real label taxonomy (supersedes the "not yet built" placeholders in file 01)

File `01-email-response-playbook.md` describes an "Admin Actions Needed" label as **not yet built**. That's outdated — the actual Gmail account already has a working `Claude/*` label hierarchy in place:

- **`Claude/Need admin (high priority)`** — email has a reply drafted (or doesn't need one) but requires backend work first (Stripe/Teachable refund or cancellation, verifying a dispute, etc.) before anything can be sent or considered resolved.
- **`Claude/Need sent`** — a customer-facing reply has been drafted and is ready to send as-is; no backend action needed, just review and hit send.
- **`Claude/Need deleted`** — no reply needed (spam, pitches, automated notifications, resolved threads, etc.); flagged for a human to confirm and delete rather than Claude deleting outright.
- **`Claude/Possible spam`** — exists, not yet used in practice; presumably for ambiguous spam calls.
- **`Claude/Possible unsubscribe`** — exists, not yet used in practice; presumably for "please stop emailing me" signals.
- **`Claude`** — parent label, not applied directly.

Routine, every time a thread is triaged: apply the correct `Claude/*` label **and** remove the `INBOX` label (via `Gmail:unlabel_thread`) so the inbox view actually clears. Applying the label alone isn't enough — Gmail keeps a thread visible in the inbox until `INBOX` is explicitly removed. This is what "empty the inbox every day" (file 00) means in practice with this label system.

If a thread already has a stale/outdated **Claude/*** label from a prior, now-resolved issue (e.g., an old admin-needed label left over after a human already handled the backend work), correct it — swap in the accurate current label rather than leaving both.

Draft-creation conventions

- **Always set **to** explicitly** to the actual customer email address (usually parsed out of the message body for contact-form submissions), rather than relying on **replyToMessageId** alone — contact-form notification emails route strangely (sender/recipient both show as **contact@speakeenglishwithvanessa.com**), and explicit **to** avoids misdirected replies.
- **Two-draft pattern for anything needing backend work:** create an internal admin-checklist draft *first* (**to:** **contact@speakeenglishwithvanessa.com** as a safe placeholder, subject  **ADMIN TASKS – DO NOT SEND**, body listing exactly what needs to happen in Stripe/Teachable/Kit), then create the customer-facing draft *second* — Gmail surfaces whichever draft was created most recently as the primary one in the thread view, so this order makes the customer draft the one that's visible first.
- **No em dashes in student-facing content** — use a comma or colon instead. (Internal notes/checklists can still use dashes.)
- **Footer format**, every draft:

Vanessa

English Teacher

www.SpeakEnglishWithVanessa.com

Hyperlinked in **htmlBody** to <https://www.SpeakEnglishWithVanessa.com>; plain text in **body** with no raw <https://> (reduces Gmail's tracking-redirect wrapping).

- **Links as hyperlinked text** ("here," "click here to join") rather than raw URLs, in both **body** (where possible) and **htmlBody** — except the Teachable sign-in link (https://speak-english-with-vanessa.teachable.com/sign_in), which stays as a plain visible URL so students can easily copy/bookmark it.
- **Sales page by default, checkout link only if asked:** per the confirmed July 27 correction in file 04, default to the FFC sales page (speakenglishwithvanessa.com/flucency/) or the Vanessa AI sales page, not the direct checkout/purchase link — only send the checkout link if the student says they can't find where to pay.

Vanessa AI sales page — gap in the knowledge base

Files 03/04/06 all reference "the Vanessa AI sales page" and instruct Claude to hyperlink it, but **no file actually contains the URL**. Found via live web search:

<https://speakenglishwithvanessa.com/speak-with-vanessa-ai/>. This should be added to [03-products-pricing-and-current-promos.md](#) under the "Speak With Vanessa AI" section so it's authoritative instead of something Claude has to re-discover each time.

Recognizing the Yearly-upgrade automated email from a reply alone

Confirmed by the person, July 28: FFC Monthly Membership students automatically receive an upsell email within their first month (per file 03's "Upsell emails" note) offering the Yearly plan at \$19/mo (5 months free vs. paying \$35/mo monthly). **The word "Upgrade" in that email's subject/body/CTA is the tell** ("Upgrade to the Yearly Membership") — if a reply is quoting or responding to an email using that language, the recipient is already an existing, paying Monthly Member, not a new lead being pitched the club for the first time.

This matters for interpreting confused replies: a student replying to this specific email with something like "I already paid \$X" is talking about their *existing* membership, not expressing confusion about whether to buy in the first place. Don't answer these the way you'd answer a new-lead inquiry — acknowledge their existing membership first, then address whatever confusion or question they raised. Case in point: riyaskismet@gmail.com replied "I already paid 45 dollars" to this exact email; the first draft answered it as if his membership status was unclear, the corrected version opens by confirming he's already a member and the offer was an upgrade, not a new purchase, before asking for the clarification needed to check the \$45 discrepancy.

Grammar accuracy — zero tolerance, not a style preference

Confirmed by the person, July 28: **grammatical accuracy in student-facing drafts is a hard requirement, not a nice-to-have**, because SEWV is a business that teaches English professionally. A single comma splice or run-on undermines credibility in a way it wouldn't for an unrelated business.

Comma splices are the specific, recurring failure mode to watch for. A comma splice is two independent clauses (each could stand alone as a full sentence) joined by only a comma, with no conjunction. This showed up repeatedly across a batch of July 28 drafts, almost always in the same shape: a warm/casual aside tacked onto a sentence with a comma instead of a period, semicolon, or conjunction. Examples of the pattern to catch before sending:

- ❌ "Thanks for the extra detail, that's really helpful." (two independent clauses, comma only)
- ✅ "Thanks for the extra detail. That's really helpful." — or add a conjunction: "Thanks for the extra detail, and that's really helpful."
- ❌ "...refunded this payment, you should see it back in your account in 5-10 business days."
- ✅ "...refunded this payment. You should see it back in your account in 5-10 business days."

Before finalizing any draft, re-read every sentence and check: could each half of a comma-joined pair stand alone as its own sentence? If yes, it's a splice, fix it with a period, semicolon, or conjunction (and/but/so/or), not a coordinating comma alone. This check applies to every draft going forward, not just ones flagged after the fact.

Stripe dispute notices — always Need deleted, not Need admin

Confirmed by the person, July 28: **Stripe dispute/chargeback notifications go straight to Claude/Need deleted, full stop.** These are generally ignored, no verification step, no admin follow-up.

This corrects an earlier interpretation this session pulled from [02-admin-task-procedures.md](#) / the playbook's dispute entry (which describes an admin follow-up of checking Teachable to confirm the student's membership was actually canceled before the dispute). Two dispute notices from July 28 (Bank Islam Malaysia \$34.56, Commonwealth Bank of Australia \$35) were initially routed to **Claude/Need admin (high priority)** with that verification note attached, then corrected to **Claude/Need deleted** per this standing rule. Going forward, skip the verification step entirely and route every Stripe dispute notification straight to **Claude/Need deleted** on sight, no exceptions, no admin note needed.

"Which module should I start with?" — near month-end rule

Confirmed by the person, July 28: when a student asks which month's module to start with (a genuinely common question near the end of a month), **if there are fewer than 5 days left in the current month, recommend starting with next month's lessons instead of the current month's.** Otherwise, point them to the 2026 Study Plan doc as usual (it names whatever the current module is). This came up with wa5aax0@gmail.com, who had 4 days left in July and was pointed to August's module instead of July's.

Bare, no-detail inquiries — treat as not-serious, delete rather than reply

Confirmed by the person, July 28: a message with essentially no content beyond something like "I need to learn English" or "I want to learn English with you," with no specifics about goals, level, timeline, or an actual question, generally signals someone who isn't a serious prospective student, not a genuine lead worth a course-recommendation reply. Default to **Claude/Need deleted** for these rather than drafting a reply, even though the general instinct elsewhere in this workflow is "when in doubt, draft a reply for anything with a real ask." A bare statement with zero personalization or specificity isn't a real ask in this sense.

This reverses an earlier call this session: khadijetou908@gmail.com's message ("I need to learn English") was initially drafted a course-recommendation reply; corrected to **Claude/Need deleted** per this rule. Contrast with genuinely low-effort-but-real questions (e.g., "is this free or paid," "which classes, beginner or advanced") — those still get a reply; the distinguishing factor is whether there's *any* specific question or detail, not just message length.

Financial hardship / can't-pay / government restrictions / internet issues — ALWAYS delete, no reply, no exceptions

Confirmed by the person, July 28, superseding everything below and the "heartfelt" exception in 01-email-response-playbook.md #19: any message that mentions lack of money, can't afford it, government restrictions, sanctions, or internet/connectivity issues preventing access or payment gets **Claude/Need deleted**, with no reply, regardless of length, tone, or how heartfelt it reads. **The "long/heartfelt hardship gets a warm reply" exception is removed entirely.** File 01's language distinguishing short can't-pay one-liners from longer heartfelt ones no longer applies, treat all of them the same way: delete, no reply.

This volume is high (per the person: "we get a ton of emails like this every day"), so the read needs to be fast and decisive, not a judgment call weighed case by case. **When in doubt about whether a message counts, err toward deleting it** — the person explicitly said to err on the side of "very very rarely" adding a warm reply to anything in this category. If a message is even partly about money, affordability, payment access, sanctions, government blocking, or internet/connectivity problems, that's the signal, delete it.

This corrects several calls made earlier in this same session, all now reversed to **Claude/Need deleted**:

- wali23426@gmail.com — couldn't afford to join
- kingshery262@gmail.com — family can't afford lessons

- noorehira255@gmail.com — described her financial situation at length (previously treated as the "genuinely substantive" contrast case that still deserved a reply, that distinction no longer applies)
- zabih.narimani20@gmail.com — sanctions preventing payment
- mah1992.01.17@gmail.com — government restrictions + internet difficulty + sanctions (already corrected once this session under a narrower rule; that narrower rule is now superseded by this blanket one)
- borisniyimenya@gmail.com — couldn't afford to join (already corrected once this session under a "short vs. heartfelt" calibration rule; that calibration distinction no longer matters, since heartfelt ones are deleted too now)

Reminder for Claude/Need admin cases: this rule is about students who can't pay / haven't paid and are explaining why, not about refund requests from students who already paid. A message like Mike Poutou's or Bhumika's ("please refund my charge") is a different category entirely (an existing paying customer requesting money back) and still gets a reply + admin follow-up as before. The line is: **already paid and wants it back → reply + admin action.** **Never paid / can't pay and explaining why → delete, no reply.**

Clarifying the line against playbook #6 (payment troubleshooting): file 01 #6 covers a student who *has* a valid card but it's failing/declining, a genuine technical problem worth troubleshooting with a reply. It does **not** cover a student who says they have no online payment method or no card at all ("I don't have any online or card payment"), that's a can't-pay case in different wording and falls under this delete-all rule, not #6. The distinguishing question: is there a specific payment method that's failing (→ #6, reply), or is there no viable payment method at all (→ delete)? Corrected example: jasminmatilda5@gmail.com (Sri Lanka, "I don't have any online or card payment") was initially treated as a #6 troubleshooting case and drafted a reply; corrected to **Claude/Need deleted**.

Recalibration: default to delete unless there's a concrete, on-topic ask

Self-assessment, confirmed by the person, July 28: roughly 8 of ~28 "Need sent" drafts from the July 28 session needed correction to "Need deleted." That's a real miscalibration, not noise, and the root cause was using "does a KB template match this message" as the signal for "does this deserve a reply," rather than "is this a serious, actionable inquiry from someone likely to be or become a paying customer, or an existing customer with an actual account matter." The knowledge base has a template for nearly everything (hardship, "are you real," bare interest, personal chit-chat) because it was written to cover the full range of things people say, not because every instance of that pattern is worth a reply given the daily volume this business receives.

Corrected default going forward: unless a message is (a) from an existing paying customer with a genuine account/billing/access issue, or (b) contains a specific, concrete, on-topic question that signals real interest or purchase intent, default to **Claude/Need deleted**. This explicitly includes, beyond the financial-hardship category already covered above:

- **Personal/conversational engagement bait** — "How are you? How did you get my email?", "are you real?", or similar messages trying to strike up personal conversation with Vanessa rather than asking a real question about the course. Delete, don't engage, even though a KB template technically covers the surface pattern ("why do I get automatic emails"). Corrected examples: hus16h@gmail.com ("How are you? How u get my email?"), mishalnydm@gmail.com ("are you real").
- **Bare "do you have anything free" asks** — a plain request for a free course/resource with no other content, personalization, or engagement. Corrected example: natnaelkibreab2001@gmail.com ("if you have for free r English coures tell me"). Initially drafted a reply pointing to the free YouTube channel; corrected to **Claude/Need deleted** — this isn't a serious lead, it's someone fishing for free stuff with zero investment in the message itself.
- Bare, no-detail inquiries (see the dedicated section above).
- Anything hardship/financial/restriction-related (see the dedicated section above).

When still uncertain, ask: would a template-matching reply actually move this person toward being a customer, or would it just be a polite response to noise? If the honest answer is "probably noise," delete it.

Vague "I need help" with no specifics — delete, don't ask a clarifying question

Confirmed by the person, July 28: when a message says something like "I need your help" with no indication of what the actual question or problem is, **don't draft a reply asking them to clarify. Just delete it.** This is a related but distinct case from the "bare, no-detail inquiries" rule above (that one's about non-serious leads with no real ask; this one's about genuine help-requests that are too vague to act on) — in both cases the resolution is the same: delete, don't engage, and don't spend a reply fishing for information the student didn't volunteer.

This also came up in a message replying to a **free-worksheet/PDF email** (part of the newsletter/lead-magnet sequence), not a course-specific email. Replies to free-resource emails are lower-signal by default than replies to actual course/account emails, worth factoring in when judging whether something's a serious lead versus noise.

Corrected example: muneerahmedmastoi3@gmail.com replied "Mam I need your help" to a free PDF worksheet confirmation email. Initially drafted a reply asking what they needed; corrected to **Claude/Need deleted**.

Always open with "Hi [Name]," — infer the name if not obviously given

Confirmed by the person, July 28: every draft should open with a personalized "Hi [Name]," greeting, not a generic "Thanks for your email!" with no name at all. Several July 28 drafts skipped the greeting entirely and launched straight into the body, that's a gap, not a stylistic choice, fix it going forward.

How to find the name, in priority order:

1. A name given in the sender's own signature (most reliable).
2. A name Vanessa's own automated email already used when addressing them (visible in the quoted original, e.g. "Hi Mohammad Saada," confirms the name to use).
3. A reasonable first-name guess parsed from the email address itself (e.g. nargisomariyan@gmail.com → "Nargis").
4. If none of the above yield anything confident, **Hi there**, is an acceptable fallback per file 00's existing guidance, don't force a guess that might be wrong.

Corrected examples from this session: mohammadccky@gmail.com (name confirmed from the quoted original, "Hi Mohammad Saada") and nargisomariyan@gmail.com (name guessed from the email address, "Nargis") both had the greeting added after initially being drafted without one. wa5aax0@gmail.com and admin@solucionesdeingenio.com had no name available anywhere, so they got **Hi there**, instead.

Country-based no-payment-access rule supersedes other reply templates

Confirmed by the person, July 28: the existing rule in file 04 (students from Bangladesh, Pakistan, and usually India generally lack payment access, so default to no paid-course offer and no reply, **Claude/Need deleted**) **takes precedence over other templates that would otherwise trigger a reply**, including the Beginner-course template (3+ sentences, mentions beginner, says something kind). If a student's location is one of these and the message doesn't fall under the "can't access the free PDF" exception, delete without a reply, even if the message would otherwise clearly qualify for a specific reply template on content alone.

Iran added to this list, confirmed July 28. The no-payment-access country list is now: **Bangladesh, Pakistan, India, and Iran.** This is on top of the general financial-hardship/sanctions/restricted-access rule elsewhere in this addendum (which already covers *explicit* mentions of sanctions or can't-pay) — this country list applies even when the message doesn't explicitly mention payment difficulty at all, just their location. Corrected example: motamedzadeh@gmail.com signed off mentioning Iran ("a warm hello from Iran"); had

no payment/hardship language in the message itself, just enthusiasm and a request for guidance on where to start. Initially drafted a course-recommendation reply; corrected to **Claude/Need deleted** per the expanded country list.

This resolves a conflict flagged earlier this session: tanseemashafi02@gmail.com (signed off "from Jammu and Kashmir," India) qualified for the Beginner template on message content (3+ sentences, mentioned being a beginner, said something kind), but per this rule her location takes precedence, corrected to **Claude/Need deleted**, no reply.

Drafts threaded off a trashed message may not surface properly

This is a common, expected pattern, not a rare edge case. Confirmed by the person, July 28: Vanessa's normal workflow is to delete (trash) an email after replying to it. This means a thread can very frequently have its earlier messages sitting in Trash, with a newer student reply arriving later on the same thread. This should be treated as the default assumption for any thread with prior back-and-forth, not something to only check occasionally.

A draft was created earlier with **replyToMessageId** pointing to a message that itself carried the **TRASH** label. The person couldn't find the resulting draft in their normal Drafts view, even though the API call succeeded and returned a draft ID.

Given how often this will come up, always check whether the specific message being replied to (via **replyToMessageId) carries a **TRASH** label before drafting, on every thread with more than one message, not just when something seems off. If it does, use the most recent non-trashed message in the thread instead.** A thread can have some messages trashed and others not; threading a new draft off a trashed message risks the draft becoming hard to find in the normal Gmail UI, even though it technically exists.

Corrected example: xicusamsj@gmail.com's Vanessa AI follow-up draft was originally threaded off **19f9941569c391e2** (carries **TRASH**), recreated threaded off **19fa92caa6f088b8** (his latest message, not trashed) instead. This exact fix was later confirmed working end-to-end in the automated routine's Newsletter run (his follow-up "I saw it, thanks" was handled correctly).

Simple, short, kind language — avoid advanced phrasing

Confirmed by the person, July 28: keep student-facing language simple and short, not just grammatically correct. Many students are English learners themselves, so idiomatic or advanced phrasing can be harder to parse than plain, direct sentences, even when it's technically fine grammar. Flagged example: "It's totally your call whether X or Y" was called out as too advanced; replaced with something more direct. When drafting, prefer short sentences and everyday words over more natural-sounding-to-a-native-speaker idioms, and generally favor brevity over thoroughness, don't over-explain every angle of a policy if a shorter answer covers what the student actually needs.

Always offer to help a student cancel directly

Confirmed by the person, July 28: whenever a reply discusses cancellation as an option (not just an active refund/cancel request already being processed), **explicitly offer to do the cancellation for them** rather than just describing that canceling is possible and leaving them to figure out how. E.g., "If you'd like me to cancel your membership for you, just let me know!" This applies to cases like Yana's (asking about pausing, cancellation mentioned as an alternative) where cancellation wasn't the direct ask but came up as an option in the reply.

Reference: how Teachable/Stripe subscription cancel + reactivate actually works

New information from Teachable support, shared by the person July 28. Useful context for understanding confusing "I canceled but got charged again" cases (a recurring theme in today's admin items), though the person notes this specific reactivation scenario is **very unusual**, not a go-to explanation to reach for by default.

- **If a student reactivates before their current billing cycle ends:** the subscription just continues as normal, no immediate charge, and they're billed as usual at the end of that cycle, exactly as if they'd never canceled.
- **If the billing cycle ends before they reactivate:** the subscription is then fully canceled, and they'd need to make a new payment to re-enroll rather than simply reactivating.
- **The mechanics:** when a student cancels, a cancellation confirmation is sent (visible on Teachable's end). Their subscription then stays active through the rest of that billing period, with no further charges scheduled. If they change their mind before the period ends, they can reactivate from the Stripe Customer Portal, and if they do, the subscription simply continues as active.
- **Practical implication:** a student could cancel and then reactivate before the period closes out, which would explain a charge that looks like "I canceled but got billed anyway" without anything actually going wrong on SEWV's end. Worth keeping in mind as a possible explanation during admin lookups on confusing billing cases, but again, this is an edge case, not the first thing to assume.

Self-sent newsletter copies — always label Need deleted, not just unread-cleared

Confirmed by the person, July 28: Vanessa is subscribed to her own newsletter, so some emails under the SEWV Newsletter label are just the broadcast copy sent from contact@speakeenglishwithvanessa.com to itself, not a real customer reply. **These should get the full Claude/Need deleted label, the same as any other no-reply item, not just have their unread status cleared without a label.** Earlier this session, two such threads

were unread-cleared without applying the delete label (an inconsistent shortcut); corrected to add the label properly. Identify these by sender = recipient = contact@speakeenglishwithvanessa.com, no third-party reply involved.

Teammate handoff labels keep UNREAD — only Claude/* labels clear it

Confirmed by the person, July 28: when routing a thread to another teammate's label (e.g., [Amber/Amber Pronunciation](#) for pronunciation-feedback audio, which Claude can't evaluate), **remove INBOX but do NOT remove UNREAD**. Teammates rely on their own unread flag to know something's waiting for them in their folder, clearing it would make the handoff invisible to the person who actually needs to act on it.

This is different from the [Claude/*](#) labels (Need sent, Need admin, Need deleted), where clearing UNREAD is correct, those represent Claude having actually resolved or drafted something, not a pending handoff to a human. The distinction: **Claude/* label applied → Claude did the work, clear UNREAD. Teammate label applied (Amber, or any other team member's folder) → handed off, not resolved, keep UNREAD.**

Corrected example: a routine run labeled Ted's pronunciation-audio thread (shinty91@naver.com) to [Amber/Amber Pronunciation](#) and cleared UNREAD along with INBOX, following the same pattern as Claude/* labels. Re-marked UNREAD afterward. Both the Inbox and Newsletter routine prompts should be updated to reflect this distinction explicitly.

Open item

The SEWV Newsletter label contains at least one confirmed pre-existing thread (July 11, 2026) mixed in among the July 28 activity that was triaged this session, and there may be more older backlog in there beyond what's been checked so far. Per the person's instruction, only *new* activity gets triaged under "check email" — any older backlog in this label (or the inbox) will be addressed separately, on request, not automatically swept up.

Action items pending — snapshot as of July 28, 2026 session

This is a point-in-time list of what's sitting in the [Claude/*](#) labels waiting on a human after the July 28 session (inbox + SEWV Newsletter, new activity only). It will go stale as items get processed — don't treat it as a live queue, check the labels directly for current state.

Claude/Need admin (high priority) — needs backend work (Stripe/Teachable/Kit) before the drafted reply can send:

1. Mike Jovane Poutou (mikejovanepoutou@gmail.com) — refund \$35 charge from July 27 + cancel subscription. Two separate contact-form threads from him, same issue, only one reply needed.
2. Joshi Bhumika (joshibhumika1998@gmail.com) — cancel subscription + refund most recent charge. Two separate messages from her (one via contact form, one direct with clearer dates: bought June 23, auto-renewed July 23), same issue, only one reply needed.
3. Roseline (roselinek123@gmail.com) — refund her original \$1 charge. Separate thread from her earlier cancellation (already done); this is just the outstanding \$1.
4. Muhammed Riyas kunnanam purth (riyaskismet@gmail.com) — replied "I already paid 45 dollars" to the automated Yearly-upgrade email (confirmed: he's an existing Monthly Member being offered the upgrade, not a new lead, see the "Recognizing the Yearly-upgrade automated email" section above). \$45 doesn't match his expected \$35/month rate. Needs a Stripe/Teachable lookup to see what he's actually being charged before we can answer him properly. Two drafts exist in this thread from an earlier, less accurate pass, the corrected customer draft and admin note (both marked "UPDATED" / dated latest) supersede them, disregard the originals.

Claude/Possible unsubscribe — explicit "stop emailing me" requests, need Kit unsubscribe action (Procedure O), no reply needed:

1. kiju4648@gmail.com
2. samsonik0000@gmail.com